

Operating System Assignment - 4

Distributed deadlock detection simulation:

fragments:

$$S_1 = P_1 \rightarrow P_2, P_3 \rightarrow P_4$$

$$S_2 = P_2 \rightarrow P_5, P_5 \rightarrow P_6$$

$$S_3 = P_6 \rightarrow P_1$$

a. Global wait-for graph (combined)

$P_1 \rightarrow P_2 \rightarrow P_5 \rightarrow P_6 \rightarrow P_1$ (cycle). Also $P_3 \rightarrow P_4$ (separate)

b. Deadlock?

→ Yes. Processes involved in deadlock cycle: P_1, P_2, P_5, P_6 .

c. Suggested distributed algorithm:

→ Use the Chandy-Misra-Harvard edge-chasing (Probe) algorithm for distributed deadlock detection - each node sends probes along wait-for edges to detect cycles without centralized graph assembly.

Distributed file system performance:

Given: local = 5 ms, remote = 25 ms, prob(remote) = 0.3

a. Expected access time:

$$E = 0.7 \times 5 \text{ ms} + 0.3 \times 25 \text{ ms} = 3.5 + 7.5 = 11 \text{ ms}$$

b. Caching Strategy:

Client-side read cache with LRU + TTL-based validation

→ Justification: frequently-read remote files will be served locally reducing remote access (0.3 fraction), LRU evicts less used items. TTL keeps staleness bounded. Improves average latency while keeping consistency manageable.

Q1.
Ans

Race condition (real-world example) and mutual exclusion.
 Two people (A and B) share a single car key in an office. Both check the key box at the same time, see the key present, and both take it - resulting in no key for either when they try to leave. This is a race condition: the outcome depends on timing of access.

Mutual exclusion solution: make the key box accessible to one person at a time (e.g. a lock). Only one person can take the key after acquiring exclusive access; others wait. This prevents conflicting.

Q2.
Ans

Peterson's solution vs semaphores (implementation complexity & hardware dependency).

Peterson's solution: simple algorithm for two processes using shared flags and a turn variable.

Semaphores: higher-level OS primitive with blocking, wakeup and queue management. More complex to implement but less hardware-dependent.

Q3.
Ans

Producers-consumers: advantage of monitors on multi-core systems.

Monitors provide automatic mutual exclusion and condition variables with blocking (no busy-waiting). In multi-core systems this reduces CPU waste (no spinning) and improved scalability by letting the OS scheduler park threads, reducing contention and cache coherence traffic compared with spin-based semaphore solutions.

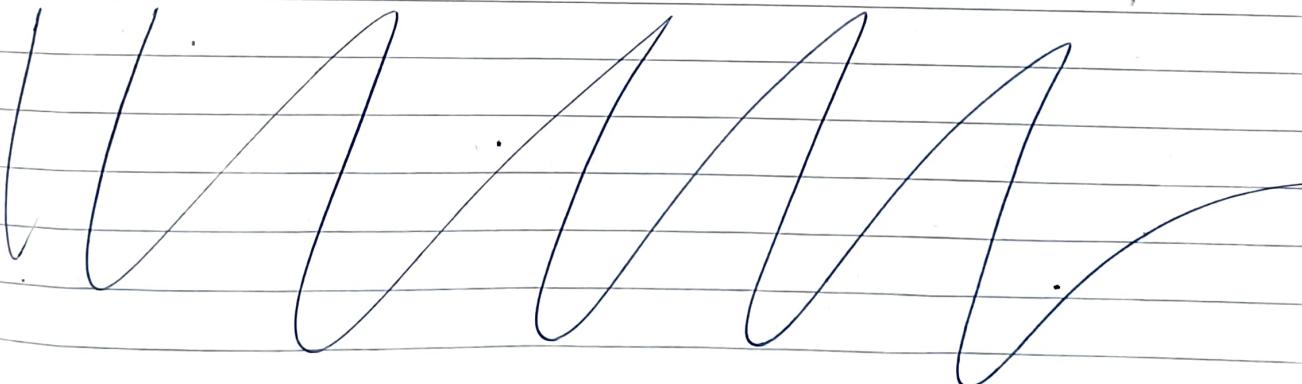
Q. 5
Ans

Reader - Writer starvation and one prevention method

Starvation Scenario: if writers are rare but continuous streams of new readers keep arriving, writer can be perpetually delayed (reader preference implementations starve writers).

Prevention: use writer preference or a fair queue (eg. queue incoming requests and access in FIFO order) or use a ticket / sequence number system so writers wait only a bounded time. This ensures fairness and prevents indefinite postponement of writers.

Q. 5
Ans Eliminating "Hold and Wait" - practical drawback.
If OS forces processes to acquire all needed resources at start (to eliminate hold-and-wait), processes must request all resources up-front or release held resources before requesting more. Practical drawback: very low resource utilization and poor concurrency - processes hold many unused resources while waiting for others, increasing blocking and reducing throughput user programs often cannot predict all future resource needs in advance.



8. Checkpointing mix to meet RPO = 1s

→ Given, full = 200 ms, incremental = 50 ms, RPO = 1s

a. Proposed Mix (over 10s) :

- Take one full checkpoint every 10s (at $t=0$ in the period)
- Take incremental checkpoint every 1s (at $t=1, 2, \dots, 9$)

$$\text{Total overhead (per 10s)} : 1 \times 200 \text{ ms} + 9 \times 50 \text{ ms} = 200 + 450 = 650$$

b. Reasoning :

- With incremental every 1s, the maximum work lost on failure $\leq 1s \rightarrow$ meets RPO.
- Full once per 10s bound recovery time (so restoring an older full & a small number of incrementals is feasible)
- This mix minimises full checkpoint cost while keeping incremental is feasible).
- This mix minimises full-checkpoint cost while keeping incremental frequency high enough to meet the RPO.

9. Case Study - Global E-commerce Platform:

a. Distributed scheduling challenges for flash sales:

- Spike, sudden spike in request
 - Geographic distribution & latency - data locality matters for latency and inventory correctness.
 - Heterogeneous nodes
 - Stratified server.
- Suggest algorithm for load balancing \Rightarrow Hybrid Approach

b. Fault-Tolerance Strategy (RTO & RPO):

- Active-active multi-region deployment: service run concurrently in multiple regions so failures is seamless ($RTO \approx \text{near-zero}$ at service level)

Data Strategy:

- Critical transactional data: Use synchronous replication within the region to guarantee consistency and low RPO; cross region replication can be asynchronous but with frequent replication to keep RPO small.