

name - yadav amit

roll no - 2301010218

course- B. Tech CSE

## Assignment- 1

1. Despite the evolution of hardware, why do modern systems still rely heavily on operating Systems?

Answer:

- Hardware provides raw capabilities; the OS provides abstraction (processes, files, devices) making hardware usable.
- Resource management: CPU scheduling, memory management, I/O coordination and device drivers.
- Security and access control (user accounts, permissions, isolation).
- Multitasking and concurrency support plus APIs for applications → portability and developer productivity.
- Performance optimizations (caching,

buffering) and fault tolerance features (process isolation, recovery).

2. OS for a wearable heart-rate monitor - which type and why?

Answer:

- A real-time embedded OS (RTOS) - e.g., FreeRTOS, Zephyr.
- Reason: deterministic, low-latency scheduling (hard/soft real-time constraints), small footprint, energy efficiency, reliable interrupt handling and simple task model (tasks + ISRs).
- Provides predictable deadlines for sampling/alerts and safe, low-power operation.

3. For a performance-critical environment which structure would you avoid (Monolithic, Layered, Microkernel), and why?

Answer:

- Avoid Microkernel for strict performance-critical cases because it incurs extra IPC and context-switching overhead between user-space services and the kernel.
- Monolithic kernels generally deliver better raw performance (fewer context switches), while

layered designs trade some performance for modularity and maintainability.

4. A developer claims OS structure doesn't matter as long as processes run. Support or refute.

Answer:

- Refute: Structure affects performance, security, robustness, maintainability, and extensibility.
- Examples: microkernels improve fault isolation but may reduce throughput; monolithic kernels are faster but risk system-wide crashes from buggy drivers. So structure matters for real-world requirements.

5. While debugging, you find a process switching error.

i) How analysing the PCB points to misinitialized registers or incorrect states?

Answer:

- The PCB saves CPU registers, program counter, stack pointer and process state.

Mismatch

between saved/restored register values (e.g.,

wrong PC or SP) in PCB indicates misinitialization or corrupted context.

- Inconsistent state fields (e.g., wrong process state: RUNNING vs READY) in PCB reveals scheduler/driver bugs.

ii) If a task is moved unexpectedly from running to waiting, what precisely does context switching involve?

Answer:

- Save current CPU state into the outgoing process's PCB (registers, PC, SP, flags).

- Update process state to WAITING and place it in appropriate wait queue.

- Scheduler picks a new process: update its PCB to RUNNING and restore its CPU state (load

registers, PC, SP).

- Update MMU/TLB as needed and perform any accounting (timestamping).

iii) OS needs to allocate I/O resources mid-execution. Which type of system call to use and

why?

Answer:

- Use asynchronous, non-blocking I/O system calls when possible: the process requests the I/O

and continues executing while the OS completes the operation (e.g., submit + callback/notification).

- Rationale: prevents stalling CPU-bound work, improves concurrency. Blocking synchronous calls

are simpler but may hurt responsiveness.

Part B - Application / Numerical

6. Context Switching Time Calculation:

Given: Save state = 2 ms, Load state = 3 ms, Scheduler overhead = 1 ms

a) Total context switching time =  $2 + 3 + 1 = 6$  ms.

b) Impact on multitasking:

- Each context switch costs CPU time (overhead) reducing usable CPU for tasks.

- High context-switch frequency (small time-slices) increases overhead and lowers throughput;

choose quantum to balance responsiveness vs overhead.

## 7. Thread Efficiency Check:

Given: Total time (single-threaded) = 40 s,  
threads per process (ideal conditions) = 2

Estimate execution time:

- In ideal perfect parallelism across 2 threads:  
 $40 / 2 = 20$  s.

Explanation:

- Multithreading allows overlapping I/O and CPU work, better CPU utilization, reduced idle time due to blocking calls, and potential parallel execution on multi-core CPUs. Overheads (synchronization, contention) reduce the ideal gain.

## 8. Given processes P<sub>1..P<sub>4</sub></sub> with bursts:

P<sub>1</sub>=5 ms, P<sub>2</sub>=3 ms, P<sub>3</sub>=8 ms, P<sub>4</sub>=6 ms

a) Gantt charts

FCFS (P<sub>1</sub>, P<sub>2</sub>, P<sub>3</sub>, P<sub>4</sub>):

10-5 15-8 18-16 116-221

P<sub>1</sub> P<sub>2</sub> P<sub>3</sub> P<sub>4</sub>

Non-preemptive SJF (order by burst): P<sub>2</sub>, P<sub>1</sub>, P<sub>4</sub>, P<sub>3</sub>:

10-3 13-8 18-14 114-221

P<sub>2</sub> P<sub>1</sub> P<sub>4</sub> P<sub>3</sub>

Round Robin (Quantum = 4 ms) - execution

order (time slices):

10|4|4-7|7-11|11-15|15-16|16-20|20-22|

P1 P2 P3 P4 P1 P3 P4

(See computed completion times below)

b) Average waiting Time (AWT) and Average Turnaround Time (ATT)

FCFS:

- waiting times: P1=0, P2=5, P3=8, P4=16
- Turnaround times: P1=5, P2=8, P3=16, P4=22
- AWT =  $(0+5+8+16)/4 = 7.25$  ms
- ATT =  $(5+8+16+22)/4 = 12.75$  ms

Non-preemptive SJF:

- Order: P2, P1, P4, P3
- waiting: P2=0, P1=3, P4=8, P3=14
- Turnaround: P2=3, P1=8, P4=14, P3=22
- AWT =  $(0+3+8+14)/4 = 6.25$  ms
- ATT =  $(3+8+14+22)/4 = 11.75$  ms

Round Robin (Quantum=4 ms):

- Completion times: P1=16, P2=7, P3=20, P4=22
- waiting = Turnaround - Burst:

$$P1 \text{ waiting} = 16 - 5 = 11$$

$$P2 \text{ waiting} = 7 - 3 = 4$$

$$P3 \text{ waiting} = 20 - 8 = 12$$

$$P4 \text{ waiting} = 22 - 6 = 16$$

- AWT =  $(11+4+12+16)/4 = 10.75$  ms
  - ATT =  $(16+7+20+22)/4 = 16.25$  ms
- c) which algorithm best balances throughput and turnaround?
- Non-preemptive SJF gives the lowest average waiting time (6.25 ms) and lowest average turnaround (11.75 ms) for this job mix, so it best balances throughput and turnaround here.

- Note: SJF requires knowledge of burst lengths and can starve long jobs.

9. i) Company migrating services to virtualized cloud:

- a) which OS architecture to select for scalability and security - and why?

Answer:

- Microkernel or minimal-kernel approach is preferable for strong security/isolation (small trusted kernel, services in user space).

Combine with a hypervisor (Type-1) for VM-level

isolation.

- Microkernel reduces attack surface and isolates drivers/services; hypervisors provide

tenant

separation and resource control.

b) How VMs aid isolation, management, resource optimization:

- Isolation: VMs keep faults within the guest boundary and separate tenants.
- Management: snapshots, live migration, cloning, easy upgrades and rollback.
- Resource optimization: overcommit, dynamic allocation, autoscaling, and monitoring.

ii) Smart home system - scheduling and IPC:

a) How OS uses scheduling and IPC to prioritize intrusion detection vs lighting:

- use priority-based preemptive scheduling and interrupts to ensure intrusion detection gets

CPU immediately.

- use IPC (message queues, event notifications) to pass sensor events to handlers; avoid busy-waiting.

- use priority inheritance to prevent priority inversion on shared resources.

b) Candidate scheduling algorithms:

- Rate Monotonic Scheduling (RMS) for periodic tasks with static priorities.
- Earliest Deadline First (EDF) for dynamic deadlines and higher utilization.

Priority-based preemptive scheduling with real-time extensions and priority inheritance is practical for mixed-criticality IoT controllers.

End of solutions.